

Lab 8

Enoch Kim

11:59PM April 29, 2021

I want to make some use of my CART package. Everyone please try to run the following:

```
if (!pacman::p_isinstalled(YARF)){
  pacman::p_install_gh("kapelner/YARF/YARFJARs", ref = "dev")
  pacman::p_install_gh("kapelner/YARF/YARF", ref = "dev", force = TRUE)
}
options(java.parameters = "-Xmx4000m")
pacman::p_load(YARF)
```

YARF can now make use of 7 cores.

For many of you it will not work. That's okay.

Throughout this part of this assignment you can use either the `tidyverse` package suite or `data.table` to answer but not base R. You can mix `data.table` with `magrittr` piping if you wish but don't go back and forth between `tbl_df`'s and `data.table` objects.

```
pacman::p_load(tidyverse, magrittr, data.table)
```

We will be using the `storms` dataset from the `dplyr` package. Filter this dataset on all storms that have no missing measurements for the two diameter variables, "ts_diameter" and "hu_diameter".

```
data(storms)

storms2 = storms %>%
  filter(!is.na(ts_diameter) & !is.na(hu_diameter) & ts_diameter > 0 & hu_diameter > 0)

storms2
```

A tibble: 1,022 x 13

	name	year	month	day	hour	lat	long	status	category	wind	pressure
	<chr>	<dbl>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<chr>	<ord>	<int>	<int>
##	1 Alex	2004	8	3	6	33	-77.4	hurricane	1	70	983
##	2 Alex	2004	8	3	12	34.2	-76.4	hurricane	2	85	974
##	3 Alex	2004	8	3	18	35.3	-75.2	hurricane	2	85	972
##	4 Alex	2004	8	4	0	36	-73.7	hurricane	1	80	974
##	5 Alex	2004	8	4	6	36.8	-72.1	hurricane	1	80	973
##	6 Alex	2004	8	4	12	37.3	-70.2	hurricane	2	85	973
##	7 Alex	2004	8	4	18	37.8	-68.3	hurricane	2	95	965
##	8 Alex	2004	8	5	0	38.5	-66	hurricane	3	105	957

```
## 9 Alex 2004 8 5 6 39.5 -63.1 hurricane 3 105 957
## 10 Alex 2004 8 5 12 40.8 -59.6 hurricane 3 100 962
## # ... with 1,012 more rows, and 2 more variables: ts_diameter <dbl>,
## # hu_diameter <dbl>
```

From this subset, create a data frame that only has storm, observation period number for each storm (i.e., 1, 2, ..., T) and the “ts_diameter” and “hu_diameter” metrics.

```
storms2 = storms2 %>%
  select(name, ts_diameter, hu_diameter) %>%
  group_by(name) %>%
  mutate(period = row_number())
```

```
storms2
```

```
## # A tibble: 1,022 x 4
## # Groups:   name [63]
##   name ts_diameter hu_diameter period
##   <chr>      <dbl>      <dbl> <int>
## 1 Alex      150.       46.0     1
## 2 Alex      150.       46.0     2
## 3 Alex      190.       57.5     3
## 4 Alex      178.       63.3     4
## 5 Alex      224.       74.8     5
## 6 Alex      224.       74.8     6
## 7 Alex      259.       74.8     7
## 8 Alex      259.       80.6     8
## 9 Alex      345.       80.6     9
## 10 Alex     437.       80.6    10
## # ... with 1,012 more rows
```

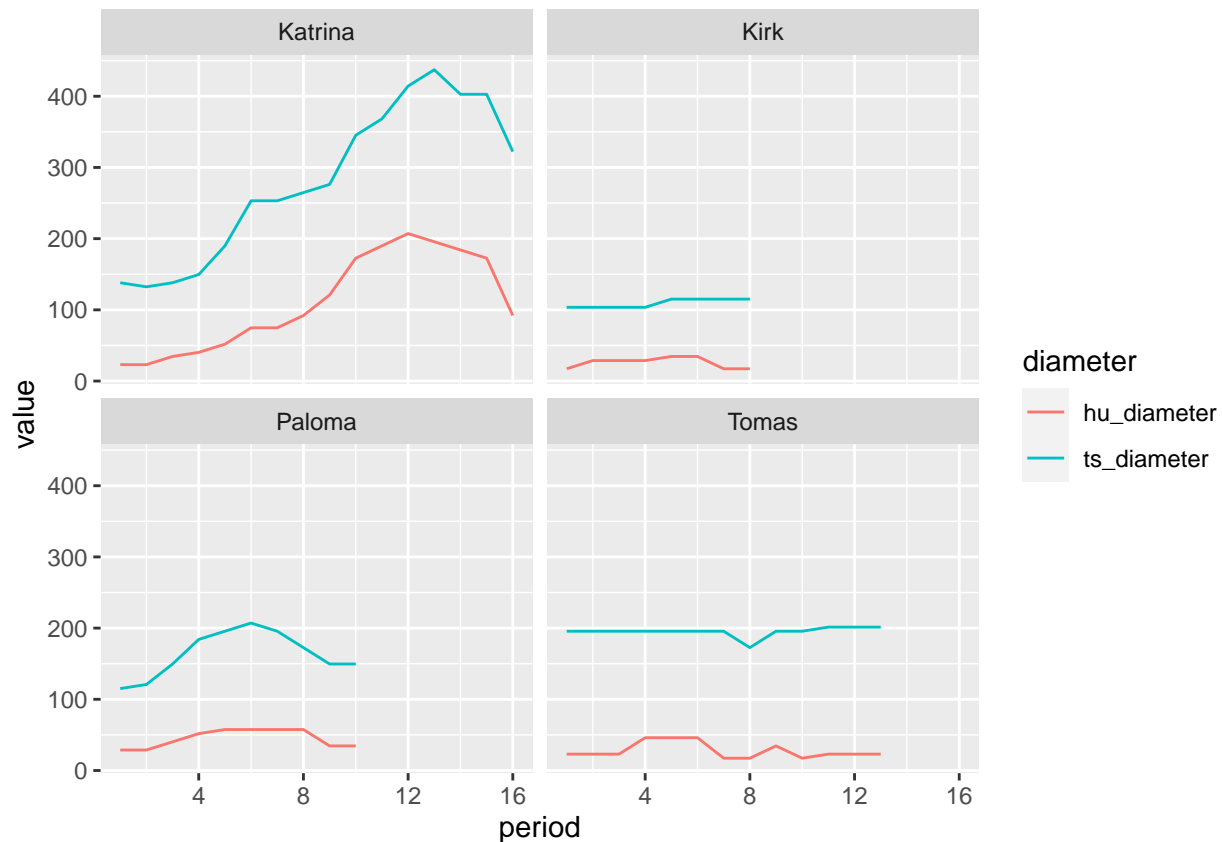
Create a data frame in long format with columns “diameter” for the measurement and “diameter_type” which will be categorical taking on the values “hu” or “ts”.

```
storms_long = pivot_longer(storms2, cols = matches("diameter"), names_to = "diameter")
storms_long
```

```
## # A tibble: 2,044 x 4
## # Groups:   name [63]
##   name period diameter value
##   <chr> <int> <chr>      <dbl>
## 1 Alex     1 ts_diameter 150.
## 2 Alex     1 hu_diameter 46.0
## 3 Alex     2 ts_diameter 150.
## 4 Alex     2 hu_diameter 46.0
## 5 Alex     3 ts_diameter 190.
## 6 Alex     3 hu_diameter 57.5
## 7 Alex     4 ts_diameter 178.
## 8 Alex     4 hu_diameter 63.3
## 9 Alex     5 ts_diameter 224.
## 10 Alex    5 hu_diameter 74.8
## # ... with 2,034 more rows
```

Using this long-formatted data frame, use a line plot to illustrate both “ts_diameter” and “hu_diameter” metrics by observation period for four random storms using a 2x2 faceting. The two diameters should appear in two different colors and there should be an appropriate legend.

```
storms_sample = sample(unique(storms2$name),4)
ggplot(storms_long %>% filter(name %in% storms_sample)) +
  geom_line(aes(x = period, y = value, col = diameter)) +
  facet_wrap(name ~. , nrow = 2)
```



In this next first part of this lab, we will be joining three datasets in an effort to make a design matrix that predicts if a bill will be paid on time. Clean up and load up the three files. Then I'll rename a few features and then we can examine the data frames:

```
rm(list = ls())
pacman::p_load(tidyverse, magrittr, data.table, R.utils)
bills = fread("https://github.com/kapelner/QC_MATH_342W_Spring_2021/raw/master/labs/bills_dataset/bills")
payments = fread("https://github.com/kapelner/QC_MATH_342W_Spring_2021/raw/master/labs/bills_dataset/payments")
discounts = fread("https://github.com/kapelner/QC_MATH_342W_Spring_2021/raw/master/labs/bills_dataset/discounts")
setnames(bills, "amount", "tot_amount")
setnames(payments, "amount", "paid_amount")
head(bills)
```

```
##           id  due_date invoice_date tot_amount customer_id discount_id
## 1: 15163811 2017-02-12  2017-01-13  99490.77   14290629   5693147
## 2: 17244832 2016-03-22  2016-02-21  99475.73   14663516   5693147
## 3: 16072776 2016-08-31  2016-07-17  99477.03   14569622   7302585
```

```
## 4: 15446684 2017-05-29 2017-05-29 99478.60 14488427 5693147
## 5: 16257142 2017-06-09 2017-05-10 99678.17 14497172 5693147
## 6: 17244880 2017-01-24 2017-01-24 99475.04 14663516 5693147
```

```
head(payments)
```

```
##           id paid_amount transaction_date bill_id
## 1: 15272980   99165.60      2017-01-16 16571185
## 2: 15246935   99148.12      2017-01-03 16660000
## 3: 16596393   99158.06      2017-06-19 16985407
## 4: 16596651   99175.03      2017-06-19 17062491
## 5: 16687702   99148.20      2017-02-15 17184583
## 6: 16593510   99153.94      2017-06-11 16686215
```

```
head(discounts)
```

```
##           id num_days pct_off days_until_discount
## 1: 5000000     20      NA              NA
## 2: 5693147     NA       2              NA
## 3: 6098612     20      NA              NA
## 4: 6386294    120      NA              NA
## 5: 6609438     NA       1              7
## 6: 6791759     31       1              NA
```

```
bills = as_tibble(bills)
payments = as_tibble(payments)
discounts = as_tibble(discounts)
```

The unit we care about is the bill. The y metric we care about will be “paid in full” which is 1 if the company paid their total amount (we will generate this y metric later).

Since this is the response, we would like to construct the very best design matrix in order to predict y.

I will create the basic steps for you guys. First, join the three datasets in an intelligent way. You will need to examine the datasets beforehand.

```
bills_with_payments = left_join(bills, payments, by = c("id" = "bill_id"))
bills_with_payments
```

```
## # A tibble: 279,118 x 9
##           id due_date invoice_date tot_amount customer_id discount_id id.y
##           <dbl> <date>   <date>         <dbl>         <int>         <dbl> <dbl>
## 1 15163811 2017-02-12 2017-01-13   99491.    14290629   5693147 14670862
## 2 17244832 2016-03-22 2016-02-21   99476.    14663516   5693147 16691206
## 3 16072776 2016-08-31 2016-07-17   99477.    14569622   7302585      NA
## 4 15446684 2017-05-29 2017-05-29   99479.    14488427   5693147 16591210
## 5 16257142 2017-06-09 2017-05-10   99678.    14497172   5693147 16538398
## 6 17244880 2017-01-24 2017-01-24   99475.    14663516   5693147 16691231
## 7 16214048 2017-03-08 2017-02-06   99475.    14679281   5693147 16845763
## 8 15579946 2016-06-13 2016-04-14   99476.    14450223   5693147 16593380
## 9 15264234 2014-06-06 2014-05-07   99480.    14532786   7708050 16957842
## 10 17031731 2017-01-12 2016-12-13   99476.    14658929   5693147      NA
## # ... with 279,108 more rows, and 2 more variables: paid_amount <dbl>,
## #   transaction_date <date>
```

```
bills_with_payments_with_discounts = left_join(bills_with_payments, discounts, by = c("discount_id" = "discount_id"))
bills_with_payments_with_discounts
```

```
## # A tibble: 279,118 x 12
##       id due_date   invoice_date tot_amount customer_id discount_id id.y
##       <dbl> <date>     <date>         <dbl>         <int>         <dbl> <dbl>
## 1 15163811 2017-02-12 2017-01-13      99491.         14290629      5693147 14670862
## 2 17244832 2016-03-22 2016-02-21      99476.         14663516      5693147 16691206
## 3 16072776 2016-08-31 2016-07-17      99477.         14569622      7302585      NA
## 4 15446684 2017-05-29 2017-05-29      99479.         14488427      5693147 16591210
## 5 16257142 2017-06-09 2017-05-10      99678.         14497172      5693147 16538398
## 6 17244880 2017-01-24 2017-01-24      99475.         14663516      5693147 16691231
## 7 16214048 2017-03-08 2017-02-06      99475.         14679281      5693147 16845763
## 8 15579946 2016-06-13 2016-04-14      99476.         14450223      5693147 16593380
## 9 15264234 2014-06-06 2014-05-07      99480.         14532786      7708050 16957842
## 10 17031731 2017-01-12 2016-12-13      99476.         14658929      5693147      NA
## # ... with 279,108 more rows, and 5 more variables: paid_amount <dbl>,
## #   transaction_date <date>, num_days <int>, pct_off <dbl>,
## #   days_until_discount <int>
```

Now create the binary response metric `paid_in_full` as the last column and create the beginnings of a design matrix `bills_data`. Ensure the unit / observation is bill i.e. each row should be one bill!

```
bills_data = bills_with_payments_with_discounts %>%
  mutate(tot_amount = if_else(is.na(pct_off), tot_amount, tot_amount*(1-pct_off/100))) %>%
  group_by(id) %>%
  mutate(sum_of_payment_amount = sum(paid_amount)) %>%
  mutate(paid_in_full = if_else(sum_of_payment_amount >= tot_amount, 1, 0, missing = 0)) %>%
  slice(1) %>%
  ungroup()
table(bills_data$paid_in_full, useNA = "always")
```

```
##
##      0      1   <NA>
## 112664 113770      0
```

How should you add features from transformations (called “featurization”)? What data type(s) should they be? Make some features below if you think of any useful ones. Name the columns appropriately so another data scientist can easily understand what information is in your variables.

```
pacman::p_load("lubridate")
bills_data = bills_data %>%
  select(-id, -id.y, -num_days, -transaction_date, -pct_off, -days_until_discount, -sum_of_payment_amount)
  mutate(num_days_to_pay = as.integer(ymd(due_date) - ymd(invoice_date))) %>%
  select(-due_date, -invoice_date) %>%
  mutate(discount_id = as.factor(discount_id)) %>%
  group_by(customer_id) %>%
  mutate(bill_num = row_number()) %>%
  ungroup() %>%
  select(-customer_id, -discount_id) %>%
  relocate(paid_in_full, .after = last_col())

bills_data
```

```
## # A tibble: 226,434 x 4
##   tot_amount num_days_to_pay bill_num paid_in_full
##   <dbl>         <int>    <int>    <dbl>
## 1    99480.           45        1        0
## 2    99529.           30        1        0
## 3    99477.           11        1        0
## 4    99479.            0        2        0
## 5    99477.           30        3        0
## 6    99477.           30        1        0
## 7    99477.            0        1        0
## 8    99477.           30        2        0
## 9    99485.           30        4        0
## 10   99477.           30        2        0
## # ... with 226,424 more rows
```

Now let's do this exercise. Let's retain 25% of our data for test.

```
K = 4
test_indices = sample(1 : nrow(bills_data), round(nrow(bills_data) / K))
train_indices = setdiff(1 : nrow(bills_data), test_indices)
bills_data_test = bills_data[test_indices, ]
bills_data_train = bills_data[train_indices, ]
```

Now try to build a classification tree model for `paid_in_full` with the features (use the `Xy` parameter in `YARF`). If you cannot get `YARF` to install, use the package `rpart` (the standard R tree package) instead. You will need to install it and read through some documentation to find the correct syntax.

Warning: this data is highly anonymized and there is likely zero signal! So don't expect to get predictive accuracy. The value of the exercise is in the practice. I think this exercise (with the joining exercise above) may be one of the most useful exercises in the entire semester.

```
#install.packages('rpart')
pacman::p_load(rpart)
mod1 = rpart(paid_in_full ~., data = bills_data_train, method = "class")
mod1
```

```
## n= 169826
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 169826 84431 1 (0.49716180 0.50283820)
##    2) tot_amount>=99291.29 63744 19157 0 (0.69946975 0.30053025)
##      4) bill_num< 1139.5 32886 2090 0 (0.93644712 0.06355288) *
##      5) bill_num>=1139.5 30858 13791 1 (0.44691814 0.55308186)
##        10) tot_amount< 99477.21 13491 4102 0 (0.69594545 0.30405455) *
##        11) tot_amount>=99477.21 17367 4402 1 (0.25346922 0.74653078) *
##    3) tot_amount< 99291.29 106082 39844 1 (0.37559624 0.62440376)
##      6) bill_num>=1237.5 20070 9933 0 (0.50508221 0.49491779)
##        12) bill_num< 3058.5 14120 5974 0 (0.57691218 0.42308782) *
##        13) bill_num>=3058.5 5950 1991 1 (0.33462185 0.66537815) *
##      7) bill_num< 1237.5 86012 29707 1 (0.34538204 0.65461796) *
```

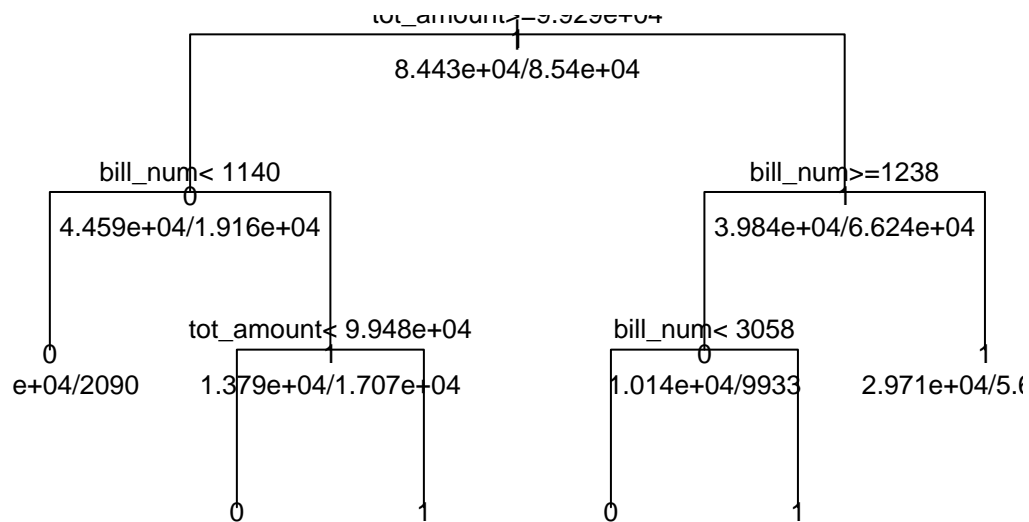
For those of you who installed `YARF`, what are the number of nodes and depth of the tree?

```
nrow(mod1$frame) ##number of nodes
```

```
## [1] 11
```

For those of you who installed YARF, print out an image of the tree.

```
plot(mod1, uniform=TRUE)
text(mod1, use.n=TRUE, all=TRUE, cex=.8)
```



Predict on the test set and compute a confusion matrix.

```
yhat = predict(mod1, bills_data_test, type = c("class"), na.action = na.pass)
oos_conf_table = table(bills_data_test$paid_in_full, yhat)
oos_conf_table
```

```
##      yhat
##      0    1
## 0 16193 12040
## 1  4099 24276
```

Report the following error metrics: misclassification error, precision, recall, F1, FDR, FOR.

```

#Levels
n = sum(oos_conf_table)
fp = oos_conf_table[1, 2]
fn = oos_conf_table[2, 1]
tp = oos_conf_table[2, 2]
tn = oos_conf_table[1, 1]
num_pred_pos = sum(oos_conf_table[, 2])
num_pred_neg = sum(oos_conf_table[, 1])
num_pos = sum(oos_conf_table[2, ])
num_neg = sum(oos_conf_table[1, ])

misclassification_error = (fn + fp) / n
cat("misclassification_error", round(misclassification_error * 100, 2), "%\n")

```

```
## misclassification_error 28.51 %
```

```

precision = tp / num_pred_pos
cat("precision", round(precision * 100, 2), "%\n")

```

```
## precision 66.85 %
```

```

recall = tp / num_pos
cat("recall", round(recall * 100, 2), "%\n")

```

```
## recall 85.55 %
```

```

false_discovery_rate = 1 - precision
cat("false_discovery_rate", round(false_discovery_rate * 100, 2), "%\n")

```

```
## false_discovery_rate 33.15 %
```

```

false_omission_rate = fn / num_pred_neg
cat("false_omission_rate", round(false_omission_rate * 100, 2), "%\n")

```

```
## false_omission_rate 20.2 %
```

```

F1 = (2 * tp) / (2 * tp + fp + fn)
cat("F1", round(F1 * 100, 2), "%\n")

```

```
## F1 75.05 %
```

Is this a good model? (yes/no and explain).

Depends, on what someone thinks a the term “good”. In this case, the best case scenario would be that FP stays lower compare to FN. When it comes to prediction, you rather have someone who is not going to pay, end up paying then having someone who was predicted to pay but not end up paying because someone paying but ends up not paying, can really damage the business. This model is “alright”

There are probability asymmetric costs to the two types of errors. Assign the costs below and calculate oos total cost.


```

C_fp = 69
C_fn = 1
cost = C_fp * fp + C_fn * fn
cost

```

```
## [1] 834859
```

We now wish to do asymmetric cost classification. Fit a logistic regression model to this data.

```

logistic_mod = glm(paid_in_full ~ ., bills_data_train, family = binomial(link = "logit"))
#p_hats_train = predict(logistic_mod, bills_data_train, type = "response")

```

Use the function from class to calculate all the error metrics for the values of the probability threshold being 0.001, 0.002, ..., 0.999 in a data frame.

```

compute_metrics_prob_classifier = function(p_hats, y_true, res = 0.001){
  #we first make the grid of all prob thresholds
  p_thresholds = seq(0 + res, 1 - res, by = res) #values of 0 or 1 are trivial

  #now we create a matrix which will house all of our results
  performance_metrics = matrix(NA, nrow = length(p_thresholds), ncol = 12)
  colnames(performance_metrics) = c(
    "p_th",
    "TN",
    "FP",
    "FN",
    "TP",
    "miscl_err",
    "precision",
    "recall",
    "FDR",
    "FPR",
    "FOR",
    "miss_rate"
  )

  #now we iterate through each p_th and calculate all metrics about the classifier and save
  n = length(y_true)
  for (i in 1 : length(p_thresholds)){
    p_th = p_thresholds[i]
    y_hats = factor(ifelse(p_hats >= p_th, 1, 0))
    confusion_table = table(
      factor(y_true, levels = c(0, 1)),
      factor(y_hats, levels = c(0, 1))
    )

    fp = confusion_table[1, 2]
    fn = confusion_table[2, 1]
    tp = confusion_table[2, 2]
    tn = confusion_table[1, 1]
    npp = sum(confusion_table[, 2])
    npn = sum(confusion_table[, 1])
  }
}

```

```

np = sum(confusion_table[2, ])
nn = sum(confusion_table[1, ])

performance_metrics[i, ] = c(
  p_th,
  tn,
  fp,
  fn,
  tp,
  (fp + fn) / n,
  tp / npp, #precision
  tp / np, #recall
  fp / npp, #false discovery rate (FDR)
  fp / nn, #false positive rate (FPR)
  fn / npn, #false omission rate (FOR)
  fn / np #miss rate
)
}

#finally return the matrix
performance_metrics
}

p_hats_train = predict(logistic_mod, bills_data_train, type = "response")
p_hats_test = predict(logistic_mod, bills_data_test, type = "response")
y_true = bills_data_train$paid_in_full
y_true_2 = bills_data_test$paid_in_full

metric_prob_classifier_in_sample = compute_metrics_prob_classifier(p_hats_train, y_true) %>% data.table
metric_prob_classifier_in_sample_tibble = as_tibble(metric_prob_classifier_in_sample)
metric_prob_classifier_out_sample = compute_metrics_prob_classifier(p_hats_test, y_true_2) %>% data.table
metric_prob_classifier_out_sample_tibble = as_tibble(metric_prob_classifier_out_sample)

```

Calculate the column total_cost and append it to this data frame.

```

C_fp = 69
C_fn = 1

metric_prob_classifier_in_sample_tibble = metric_prob_classifier_in_sample_tibble %>%
  mutate(total_cost = C_fp * FP + C_fn * FN)
metric_prob_classifier_in_sample_tibble

```

```
## # A tibble: 999 x 13
```

	p_th	TN	FP	FN	TP	misl_err	precision	recall	FDR	FPR	FOR
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	0.001	0	84431	0	85395	0.497	0.503	1	0.497	1	NaN
## 2	0.002	0	84431	0	85395	0.497	0.503	1	0.497	1	NaN
## 3	0.003	0	84431	0	85395	0.497	0.503	1	0.497	1	NaN
## 4	0.004	0	84431	0	85395	0.497	0.503	1	0.497	1	NaN
## 5	0.005	0	84431	0	85395	0.497	0.503	1	0.497	1	NaN
## 6	0.006	0	84431	0	85395	0.497	0.503	1	0.497	1	NaN
## 7	0.007	0	84431	0	85395	0.497	0.503	1	0.497	1	NaN
## 8	0.008	0	84431	0	85395	0.497	0.503	1	0.497	1	NaN

```
## 9 0.009      0 84431      0 85395      0.497      0.503      1 0.497      1  NaN
## 10 0.01      0 84431      0 85395      0.497      0.503      1 0.497      1  NaN
## # ... with 989 more rows, and 2 more variables: miss_rate <dbl>,
## #   total_cost <dbl>
```

```
metric_prob_classifier_out_sample_tibble = metric_prob_classifier_out_sample_tibble %>%
  mutate(total_cost = C_fp * FP + C_fn * FN)
metric_prob_classifier_out_sample_tibble
```

```
## # A tibble: 999 x 13
##   p_th    TN    FP    FN    TP miscl_err precision recall   FDR   FPR   FOR
##   <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.001      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## 2 0.002      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## 3 0.003      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## 4 0.004      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## 5 0.005      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## 6 0.006      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## 7 0.007      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## 8 0.008      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## 9 0.009      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## 10 0.01      0 28233      0 28375    0.499    0.501      1 0.499      1  NaN
## # ... with 989 more rows, and 2 more variables: miss_rate <dbl>,
## #   total_cost <dbl>
```

Which is the winning probability threshold value and the total cost at that threshold?

```
winning_prob_threshold_insample = which.min(metric_prob_classifier_in_sample_tibble$total_cost)
winning_prob_threshold_insample_metrics = metric_prob_classifier_in_sample_tibble[winning_prob_threshold_insample,]
cat("The winning probability threshold value in-sample is:", min(winning_prob_threshold_insample_metrics$total_cost))
```

```
## The winning probability threshold value in-sample is: 85395
```


```
winning_prob_threshold_outsample = which.min(metric_prob_classifier_out_sample_tibble$total_cost)
winning_prob_threshold_outsample_metrics = metric_prob_classifier_out_sample_tibble[winning_prob_threshold_outsample,]
cat("\n \nThe winning probability threshold value out-sample is:", min(winning_prob_threshold_outsample_metrics$total_cost))
```

```
##
##
## The winning probability threshold value out-sample is: 28375
```

Plot an ROC curve and interpret.

```
pacman::p_load(ggplot2)
metrics_in_and_out_performance = rbind(
  cbind(metric_prob_classifier_in_sample_tibble, data.table(sample = "in")),
  cbind(metric_prob_classifier_out_sample_tibble, data.table(sample = "out"))
)
ggplot(metrics_in_and_out_performance) +
```

```
geom_line(aes(x = FPR, y = recall)) +
geom_abline(intercept = 0, slope = 1, col = "Red") +
coord_fixed() + xlim(0, 1) + ylim(0, 1)
```



ROC stands for the “receiver operator curve”. The predictive power can be calculated by the area under the curve and the Red line can be used to compare probability of the models of estimation.

Calculate AUC and interpret.

```
pacman::p_load(pracma)
auc_in_sample = -trapz(metric_prob_classifier_in_sample_tibble$FPR, metric_prob_classifier_in_sample_tibble$recall)
cat("AUC in-sample: ", auc_in_sample)

## AUC in-sample: 0.6018531

auc_oos = -trapz(metric_prob_classifier_out_sample_tibble$FPR, metric_prob_classifier_out_sample_tibble$recall)
cat("\n\nAUC out-sample: ", auc_oos)


##
##
## AUC out-sample: 0.6074544
```

The AUC in-sample is 0.5960914 and the AUC out-sample is 0.5960528. Due AUC being greater than 0.5, this model has predictive power.

Plot a DET curve and interpret.

```
ggplot(metrics_in_and_out_performance) +
  geom_line(aes(x = FDR, y = miss_rate, col = sample)) +
  coord_fixed() + xlim(0, 1) + ylim(0, 1)
```

Warning: Removed 602 row(s) containing missing values (geom_path).



In this plot, the DET curve shows a trade off between FDR and FOR, I don't really know how to interpret this but when FDR is around 0.42 and the miss_rate is around .98, that seems to be the start of the optimal trade off points/plot.